

Final Project: Value-Iteration Method for MDP

Group 6

The Chinese University of Hong Kong, Shenzhen

Chuqiao Feng 120090272

Lang Tian 120090793

Linli Zhou 120090555

Manning Wu 120090753

April 21, 2023

Overview

- 1 Problems
- 2 Preliminaries
- 3 Algorithm
 - Random VI
 - CyclicVI
- 4 Numerical Experiments and Results
 - Maze Setting and Notations
- 5 Stochastic Approximation
- 6 Tic-Tac-Toe Game
 - MDP with Value Iteration
 - Q-learning Method
- 7 Summary

Preview

- We work on implementing the value-iteration method for Markov Decision Processes.
- We ask the following question:
Q: How to solve Markov Decision Process more efficiently?
- Our answers:(Overview)
A1: Value Iteration with random and cyclic strategy
A2: Value Iteration with state aggregation

Motivation

- We want to optimize sequential decision-making in reality.
- Why is the topic challenging?
 1. Curse of dimensionality
 2. The trade-off between efficiency and performance.

Problem Setup

Markov Decision Processes (MDPs)

In MDPs, we consider the following minimization problem

$$\min_{\pi \in \mathcal{A}^S} V_{\pi}(i) := \mathbb{E}_{\pi} \left[\sum_{t=0}^{\infty} \gamma^t c_{a_t}(i_t) \mid i_0 = i \right] \quad (1)$$

where $\{i_0, a_0, i_1, a_1, \dots, i_t, a_t, \dots\}$ are state-action transitions generated by the MDP under the fixed policy π , i.e. $a_t = \pi_{i_t}$, and the expectation $\mathbb{E}_{\pi}[\cdot]$ is over the set of (i_t, a_t) trajectories.

Notation

- S is finite state space, and the number of the total states is $|S|$
- A is finite state action space
- $\gamma \in [0, 1)$ is the discounted factor
- P is the collection of state-action-state transition probabilities, with $P(i'|i, a)$ represents the probability of going to state i' from state i when taking action a
- c is the collection of costs at different state-action pairs, i.e. we cost $c_a(i)$ if we are currently in state i and take action a

Bellman optimality equations

Bellman optimality equations (Bellman 1957)

The (optimal) value function achieved by the optimal policy satisfies

$$V^*(i) = \min_{a \in A_i} \left(c_a(i) + \gamma \sum_{i' \in S} P(i'|i, a) V^*(i') \right) \quad (2)$$

Value operator

For a given MDP, the value operator $T : \mathbb{R}^{|S|} \rightarrow \mathbb{R}^{|S|}$ is defined for all $U \in \mathbb{R}^{|S|}$ and $i \in S$ by

$$T(U)_i = \min_{a \in A_i} \left(c_a(i) + \gamma \sum_{i' \in S} P(i'|i, a) U(i') \right) \quad (3)$$

V^* is the fixed point of T .

Connection with Linear Programming(LP)

The Value operator inspires us to find an upper bound, $V \geq T(V)$, by considering a larger set of linear constraints.

$$V(i) \leq c_a(i) + \gamma \sum_{i' \in S} P(i'|i, a) V(i') \quad \forall a \in A, i \in S \quad (4)$$

Thus we can reformulate (1) as follows:

$$\begin{aligned} \min \quad & \sum_{i \in S} V(i) \\ \text{s.t.} \quad & V(i) \leq c_a(i) + \gamma \sum_{i' \in S} P(i'|i, a) V(i') \quad \forall a \in A, i \in S \end{aligned} \quad (5)$$

Value Iteration

Lemma (Contraction Mapping)

For all values $U, V \in \mathbb{R}^{|S|}$ we have that $\|T(U) - T(V)\|_\infty \leq \gamma \|U - V\|_\infty$ and consequently $\|T(U) - V^\|_\infty \leq \gamma \|U - V^*\|_\infty$, where V^* is the optimal value vector.*

- Based on Lemma 1, we can use the fixed-point iteration method, which is called **Value Iteration** in MDPs

$$V^{k+1} = TV^k$$

- We have $\|V^k - V^*\|_\infty \leq \gamma^k \|V^0 - V^*\|_\infty$

Lemma (Entry-wise Monotone Property)

if values $U, V \in \mathbb{R}^{|S|}$ satisfy $U \leq V$ entry-wise, then $T(U) \leq T(V)$ entry-wise.

(Vanilla) Value Iteration Algorithm

Algorithm 1 vanilla value iteration(S, A, c, P, γ)

Initialize $V^0 \in \mathbb{R}^{|S|}$ arbitrarily;

for $k = 0, 1, \dots$ **do**

for $i \in S$ **do**

$V^{k+1}(i) = \min_{a \in A_i} \{c_a(i) + \gamma \sum_{i' \in S} P(i'|i, a) V^k(i')\}$

end

end

Random Value Iteration

Motivation: When $|S|$ is huge, update is computationally expensive in each iteration.

Solution: Randomly select a subset of S to update.

Random Value Iteration (Random VI)

In the k th iteration, randomly select a subset of states B^k and do

$$V^{k+1}(i) = \min_{a \in \mathcal{A}_i} \{c_a + \gamma \sum_{i'} P(i'|i, a) V^k(i')\}, \quad \forall i \in B^k. \quad (6)$$

Modification (Influence Tree): B^k is a random subset of states which are connected by any state in B^{k-1}

Cyclic Value Iteration

Cyclic Value Iteration (CyclicVI)

Update one state at a time in order. In the k th iteration do

- Initialize $\tilde{V}^k = V^k$.
- For $i = 1$ to $|S|$

$$\tilde{V}^k(i) = \min_{a \in \mathcal{A}_i} \{c_a(i) + \gamma \sum_{i'} P(i'|i, a) \tilde{V}^k(i')\} \quad (7)$$

- $V^{k+1} = \tilde{V}^k$.

Difference with Vanilla VI: vanilla VI is synchronous (using V^k in update), but CyclicVI is **asynchronous**.

Modification (Randomly Permuted CyclicVI):

Update one state at a time in random order B^k during the k th iteration.

Maze Setting and Notations

Maze Setting

- Standard Maze: 2D Maze with m^2 states ($m = \text{height} = \text{width}$)
- Terrain Maze: 3D Maze with m^2 states
(Height: Initially, consider the height of each state as a random number. Then, the height of each state is updated by calculating the average of the heights around it.)

Example




Figure: Standard Maze




Figure: Terrain Maze

Comparison of methods in standard maze

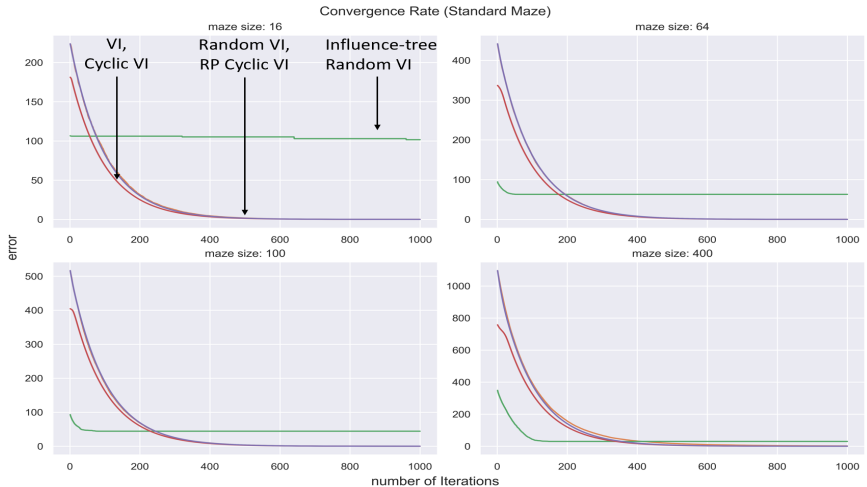


Figure: Convergence rates of various value iteration methods in standard maze.

($\gamma = 0.99$, 20000 runs, random sample size/maze size = 1)

Comparison of methods in terrain maze

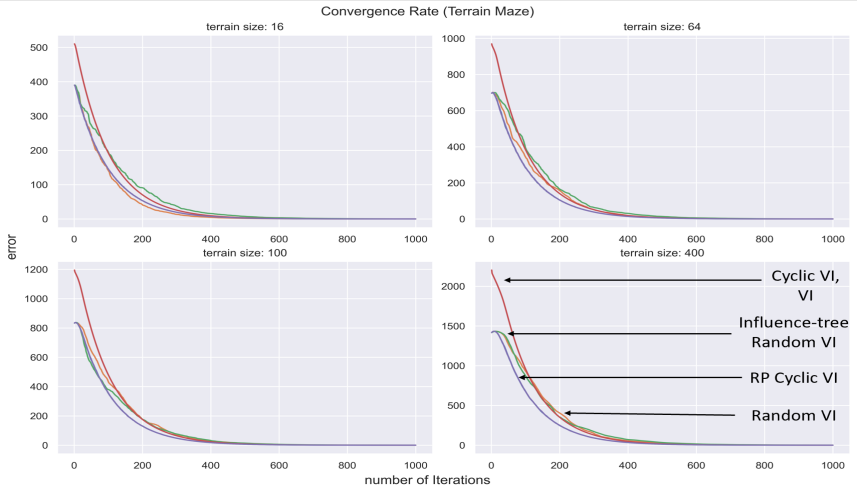


Figure: Convergence rates of various value iteration methods in terrain maze.

($\gamma = 0.99$, 20000 runs, random sample size/maze size = 1)

Results

- For both standard maze and terrain maze
 - Number of iterations to converge:
 $RP \text{ Cyclic VI} = VI < \text{Cyclic VI} < \text{Random VI} < \text{Influence-tree Random VI}$
 - Random VI and Influence-tree random VI have lower computational cost.

More computation-efficient method?

- Difficulty: The state space S is huge.
- Random selection can partially reduce the computation cost, but does not utilize the problem structure information.
- New idea: state aggregation.
 - Similar states have close values (long-term rewards).
 - We can group/aggregate such states into a mega-state.
 - The state space size is reduced.
 - Mathematical meaning: piece-wise constant function approximation rather than discrete table.

State Aggregation

Mega State

A state partition $\{S_i\}_{i=1}^K$ on \mathcal{S} : $\mathcal{S} = \cup_{i=1}^K S_i$ and $S_i \cap S_j = \emptyset$ for $i \neq j$
 Denote $W \in \mathbf{R}^K$ the cost-to-go value function for the mega state

The current value of W induces a value function $\tilde{V}(W) \in \mathbf{R}^{|\mathcal{S}|}$ on the original state space:

$$\tilde{V}(s, W) = W(j), \quad \text{for } s \in S_j$$

Adaptive Aggregation

Challenge:

- If aggregation rule W is pre-specified, we hope (and could design W such) that $\|\tilde{V}(W) - V^*\|_\infty$ is small.
- But, we do not know V^* before problem-solving, so it is hard to design W .

Solution in [Chen et al., 2021]:

- We have $V^k \rightarrow V^*$, so V^k is a surrogate of V^* .
- We can adaptively update the aggregation rule W based on V^k .

Value-based Aggregation

Algorithm 2 Value-based Aggregation

Input: $\varepsilon, \mathbf{V} = (V(1), \dots, V(|\mathcal{S}|))^T$

$b_1 = \min_{s \in |\mathcal{S}|} V(s), b_2 = \max_{s \in |\mathcal{S}|} V(s), \Delta = (b_2 - b_1)/\varepsilon;$

for $i = 1, \dots, \lceil \Delta \rceil$ **do**

$\hat{S}_i = \{s | V(s) \in [b_1 + (i-1)\varepsilon, b_1 + i\varepsilon)\}, \hat{W}(i) = b_1 + (i - \frac{1}{2})\varepsilon$

Output: Return $\{S_i\}_{i=1}^K$ and W

Key idea: Discretize V^k into intervals based on $\min_s V^k(s)$ and $\max_s V^k(s)$,

Periodical Implementation

Two-Phases Algorithms:

- Phase 1 (with \mathcal{B}): algorithm performs global updates on $|S|$.
- Phase 2 (with \mathcal{A}): algorithm performs state-aggregated updates.

For a pre-specified number of iterations n , the time horizon $[1, n)$ is divided into intervals of the form $\mathcal{B}_1, \mathcal{A}_1, \mathcal{B}_2, \mathcal{A}_2, \dots$.

- Example:

$$\begin{aligned}\mathcal{B}_1 &= \{1, 2, 3, 4\}, & \mathcal{A}_1 &= \{5\}, \\ \mathcal{B}_2 &= \{6, 7, 8, 9\}, & \mathcal{A}_2 &= \{10\}\end{aligned}$$

Algorithm

Algorithm 3 Value Iteration with Adaptive Aggregation

Input: $P, c, \gamma, \varepsilon, \{\alpha_t\}_{t=1}^{\infty}, \{\mathcal{A}_i\}_{i=1}^{\infty}, \{\mathcal{B}_i\}_{i=1}^{\infty}$

Initialize $W_0 = 0, V_1 = 0, t_{sa} = 1$

for $i = 1, \dots, n$ **do**

if $t \in \mathcal{B}_i$ **then**

if $t = \min\{\mathcal{B}_i\}$ **then**

$V_{t-1} = \tilde{V}(W_{t-1})$

for $j = 1, \dots, |\mathcal{S}|$ **do**

$V_t(j) = T_j V_{t-1}$

Continued

```

if  $t \in \mathcal{A}_i$  then
  if  $t = \min\{A_i\}$  then
    Define  $\{S_i\}_{i=1}^K$  and  $W_t$  to be the output of Algorithm 2
  for  $j = 1, \dots, K$  do
    Sample state  $s$  uniformly from collection  $S_j$ .
    
$$W_{t+1}(i) = (1 - \alpha_t)W_t(i) + \alpha_t T_i \tilde{V}(W)$$

   $t_{sa} = t_{sa} + 1$ 
  —
if  $n \in \mathcal{B}_i$  then
  return  $V_n$ 
return  $\tilde{V}(W_n)$ 

```

Experiments

Setting

- Discount factor γ : 0.99
- $|\mathcal{A}_i| : 2 \quad |\mathcal{B}_i| : 5$
- Learning rate α_t : $\frac{1}{\sqrt{t}}$
- ε : 0.5
- Initialization V_0 : $\mathbf{0}$

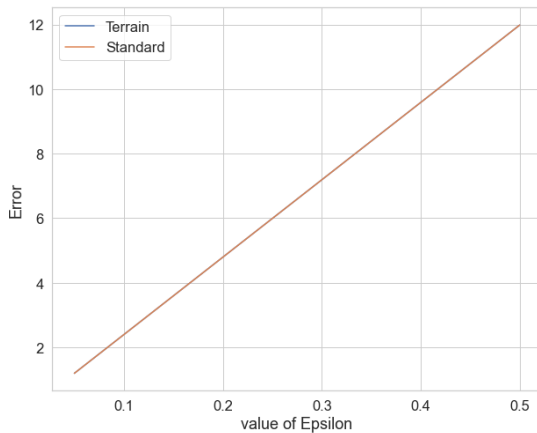
Experiments

Influence of ε : We run experiments on a 20×20 maze with different setting of ε to test the effect of ε on error.

Convergence: We test the convergence of algorithm 3 against value iteration (VI) on 20×20 standard and terrain maze.

Efficiency: We compare the computation time of algorithm3 in 4000 runs against VI on large-scale terrain maze (50×50) repeated for 20 times.

Result I

Figure: Influence of ε

$\varepsilon = 0.05, 0.2, \text{ and } 0.5$. The error $\|E_t\|_\infty \propto \varepsilon$.

Result II

orange line: Algorithm 3 **blue line:** Value Iteration

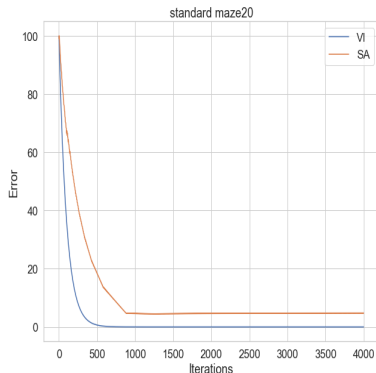


Figure: Standard maze

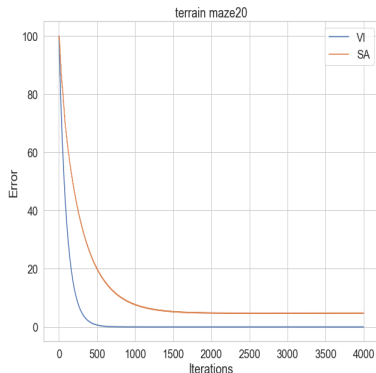



Figure: Terrain maze

The state-aggregated update \mathcal{B} will increase the $\|E_t\|_\infty$.

Result III



Scale	Time (s)	SA : VI
100	6.59 : 5.22	
400	16.35 : 21.51	
900	26.69 : 69.91	
1600	47.68 : 181.70	
2500	112.72 : 621.45	

Table: Computation Time

Figure: Efficiency test

The algorithm 3 improved efficiency in terms of computation time compared to Value Iteration.

Tic-Tac-Toe Overview

- Value Iteration
- Q-learning (Stochastic Approximation)
- Deep Reinforcement Learning (Stochastic Approximation + Function Approximation)

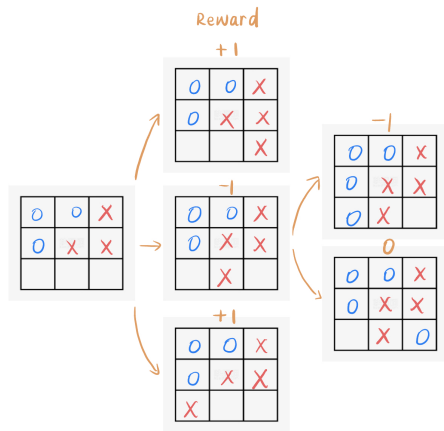


Figure: Example of rewarding state in the learning process.

Game Setting and Notations

Example: (part of a game trajectory)

States trajectory: $\{[-1, 0, 1, 0, -1, 0, 1, -1, 0], [-1, 0, 1, 0, -1, 0, 1, -1, 0]\}$

Values for X player: $\{0.9, 0.81\}$

(The current condition benefits player X)

Values for O player: $\{0.1, 0.18\}$

(The current condition is not beneficial for player O)

Choose Action: pick the position with the highest value in available space.

Agent's Value Function

38.2%	44.8%	44.8%
44.7%	98.8%	63.8%
40.7%	49.4%	50.6%

Agent's Value Function

0	89.2%	43.8%
44.2%	0	50.4%
86.9%	28.1%	61.1%

Figure: Choosing Action based on state-value function.

Value Iteration Method

- Retrieve the initialized state-value pairs (e.g. for player X)

$$\text{Value} = \begin{cases} 1, & \text{if X wins} \\ 0, & \text{if O wins} \\ 0.5, & \text{otherwise} \end{cases}$$

- Apply Value Iteration over **all the states** several times
- Go until convergence (usually not more than 3 loops)

Value-Iteration

- $V_{k+1}(s) = \max_a (R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V_k(s'))$
- Same reason as the previous problem on Bellman Equation(DP)

To retrieve the optimal policy after the value iteration:

- $\pi(s) = \operatorname{argmax}_a R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V_{k+1}(s')$

Value Iteration Method




Figure: Convergence of value iteration

Problems on Time & Space Complexity




Figure: Cost of Memory Space on 3*3 game




Figure: Cost of Memory Space on 4*4 game

Q-learning Algorithm

Algorithm 4 Epsilon-Greedy Q-Learning Algorithm

Input: α : *learning rate*, γ : *discount factor*, ϵ : *a small number*

Result: A Q-table containing $Q(S, A)$ pair defining estimated optimal policy π^*

Initialize $Q(S, A)$ arbitrarily, except $Q(\text{terminal}, .)$;

$Q(\text{terminal}, .) \leftarrow 0$

for *each episode* **do**

 Initialize state S ;

for *each step in episode* **do**

$A \leftarrow \text{SELECT} - \text{ACTION}(Q, S, \epsilon)$;

 Take action A , then observe reward R and next state S' ;

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$;

$S \leftarrow S'$

Q-learning Method

- $$Q(S, a) \leftarrow Q(S, a) + \alpha[R_a(S, S') + \gamma \max_{a'} Q(S', a')]$$







Figure: Experiment Result of Q-learning

Deep Reinforcement Learning Approach: idea

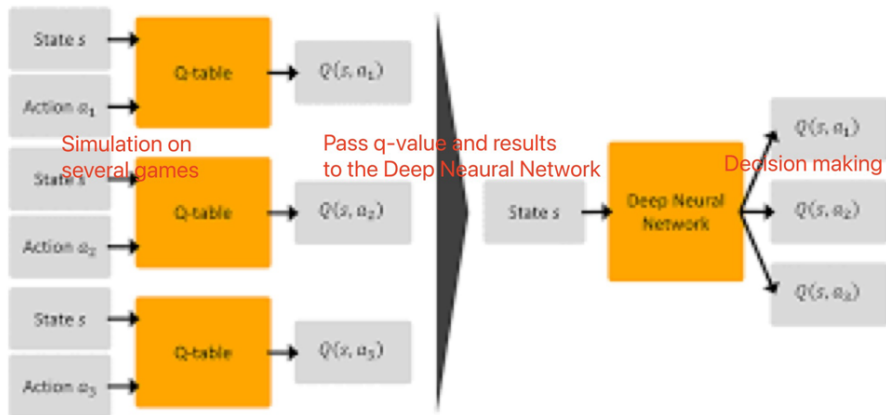


Figure: Logic of Deep Reinforcement Learning in Tic-Tac-Toe

Further: Deep Learning Approach

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 200)	2000
dropout (Dropout)	(None, 200)	0
dense_1 (Dense)	(None, 125)	25125
dense_2 (Dense)	(None, 75)	9450
dropout_1 (Dropout)	(None, 75)	0
dense_3 (Dense)	(None, 25)	1900
dense_4 (Dense)	(None, 3)	78

Figure: NN structure




Figure: NN structure

Experiment: Deep Learning Approach

Results for player 1:

Wins: 976 (97.6%)

Loss: 0 (0.0%)

Draw: 24 (2.4%)

Results for player 2:

Wins: 735 (73.5%)

Loss: 45 (4.5%)

Draw: 220 (22.0%)

Results for player 1:

Wins: 294 (29.4%)

Loss: 323 (32.3%)

Draw: 383 (38.3%)

Results for player 2:

Wins: 323 (32.3%)

Loss: 294 (29.4%)

Draw: 383 (38.3%)

Figure: Trained X v.s.
Random O

Figure: Random X v.s.
Trained O

Figure: Trained X v.s.
Trained O

Notice:

- Can not perform as well as Value Iteration in this tic-tac-toe condition.

Possible reasons:

- Total space is not large enough for deep learning
- Neural Network has low explainability for this model

Summary

- We discussed the value-iteration problem and its variance on different conditions
- Key take take-away message:
 - VI with Adaptive Aggregation shows improved efficiency in terms of computation time than basic VI.
 - Random Permuted Cyclic Value Iteration, potentially leads to further improvement in convergence speed. Because it explores the development of heuristic algorithms that can predict permutation.
 - For tic-tac-toe: Q-learning is better than Value Iteration in time and space complexity, DRL only performs well in a more complicated problem
- Potential future directions and limitations: State Aggregation only shows its efficiency in large-scale problems, we can try to find a way to solve this in a small-scale problem.

References



Chen, G., Gaebler, J. D., Peng, M., Sun, C., and Ye, Y. (2021).
An adaptive state aggregation algorithm for markov decision
processes.

arXiv preprint arXiv:2107.11053.

Acknowledgement

Thanks to the instructions from Ziniu Li

Thank you!